

Unit-II Transactions

Topics to Cover:

1) Transaction Concepts

- 2) ACID properties
- 3) Serializable and Immerse update
- 4) Shadow logging
- 5) ARIES Algorithm
- 6) Serializability
- 7) Transaction Support in SQL
- 8) Need for Concurrency
- 9) Concurrency Control
- 10) Two-phase locking
- 11) Timestamp
- 12) Multiversion
- 13) Multiple Granularity locking
- 14) Read lock Handling
- 15) Recovery concepts
- 16) Recovery based on log

Transaction Concepts

- * They are set of logically related operations
 - * They are a group of tasks
 - * A transaction is an action or series of actions
 - * It is performed by a single user
- Ex: sending 800 from A to B

open-account(A) open-account(B)

old-bal = A-bal \Rightarrow old-bal = B-bal

new-bal = old-bal - 800 new-bal = old-bal + 800

A-bal = new-bal B-bal = new-bal

close-account(A) close-account(B)

Operations of Transaction:

Read: It is used to read the value of A from the database and store them in buffer

Write: It writes the value to DB from buffer

Commit: permanently changed

Rollback: Undo changes by going to save point

ACID Properties

1) Atomicity: All or never, All operations performed either once or never

Abort: If a transaction aborts all changes are not visible

Commit: All changes are changed permanently

Example: If A=100 and B=100, if A is updated to 200 and B is not updated to 200

A=100 B=100
written written

2) Consistency: The updation must be consistent to avoid inconsistency

If money is sent the updation like - from A and B to b must be consistent.

3) Isolation: The data that is used by a transaction can not be used by other until that transaction is completed, until then the transaction has to wait.

4) Durability:

- * The durability property is used to indicate the performance of the database's consistent state. It states that the transaction make the permanent changes of these cannot be lost by the errors. When a transaction is completed then the database reaches a state known as the consistent state.
- * This state can not be lost.

States of Transaction: A transaction has at least one of the following states:

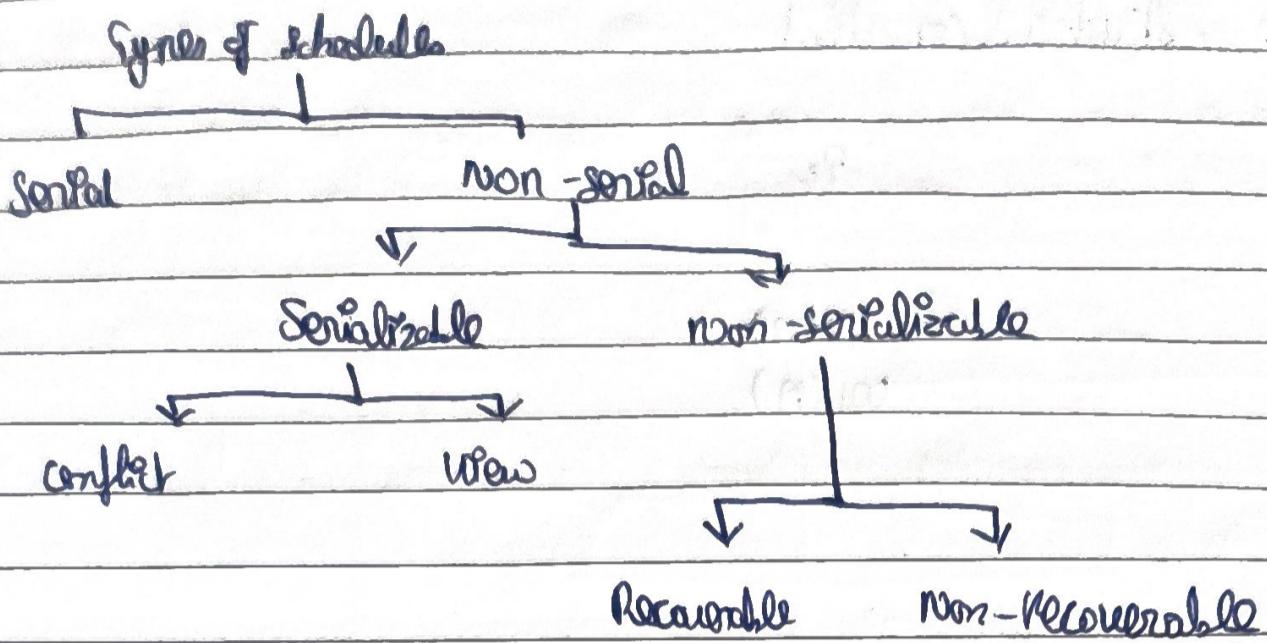
- 1) Active state
- 2) partially committed state
- 3) committed state
- 4) failed state
- 5) Aborted state

Schedule: In DBMS we have to maintain the consistency of data.

- * A transaction is a set of actions.
- * When multiple transaction has to run concurrently, the sequence to execute them they are executed one after one, this is called Schedule.

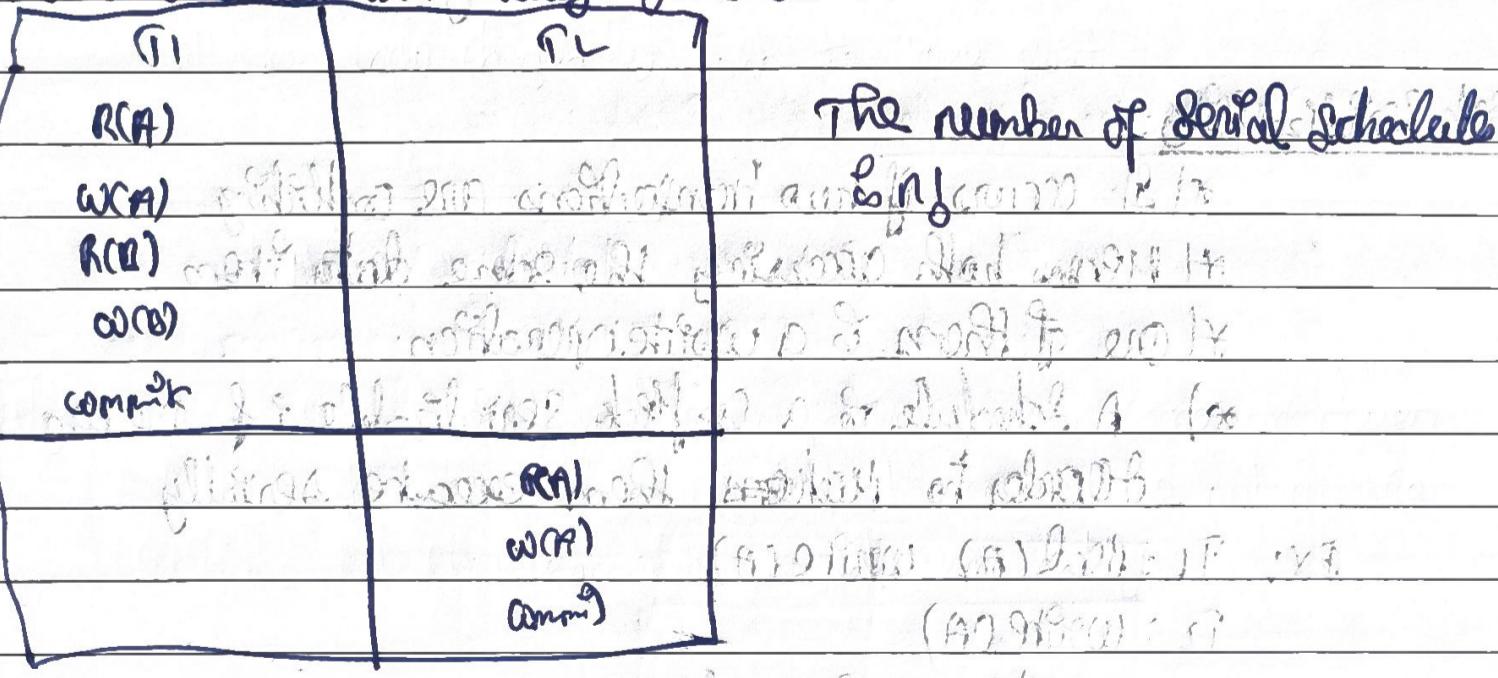
conflict-free sequence of transaction of the database will be called serializability.

* Scheduling is used to avoid concurrency problems.



Serial:

- * AS the name says each transaction is executed sequentially in serial order without affecting other, it is also called non-interleaved execution.
- * Serial schedules are always recoverable, fair, safe, strict and consistent. A serial schedule will always give correct results.



Non-Serial Schedule:

- * Multiple transactions are executed concurrently without waiting for the other.

- * Transactions proceed without completion of other
- * If either interleaved or mixed, they are not recoverable
- (Cascades, Strict and Consistent)

T_1
 R(A)
 w2(A)
 R3(A)
 w3(A)

T_2
 r2(A)
 w1(A)

Serializability

Ensures fairness of execution and consistency of data

Implementation: from bottom up at T1, with addition of locks with time

* Ensures transactions are executed serially

* If has hysteresis \rightarrow conflict serializability can't be broken
 \rightarrow View Serializability (in database block)

Conflict Serializability

* It occurs if two transactions are conflicting

* Means both accessing the same data item

* One of them is a write operation

* A schedule is a conflict serializable if we can
 & order to make them execute serially

Ex: T_1 : Read(A) Write(A)

T_2 : Write(A)

Operation	Transaction
Read(A)	T_1
Write(A)	T_2

Write(A) \rightarrow T_2 writing and T_1 has written

This is conflicting

T₁ T₂
 Read(A)
 write(A) read(A)
 write(A)
 Read(A)

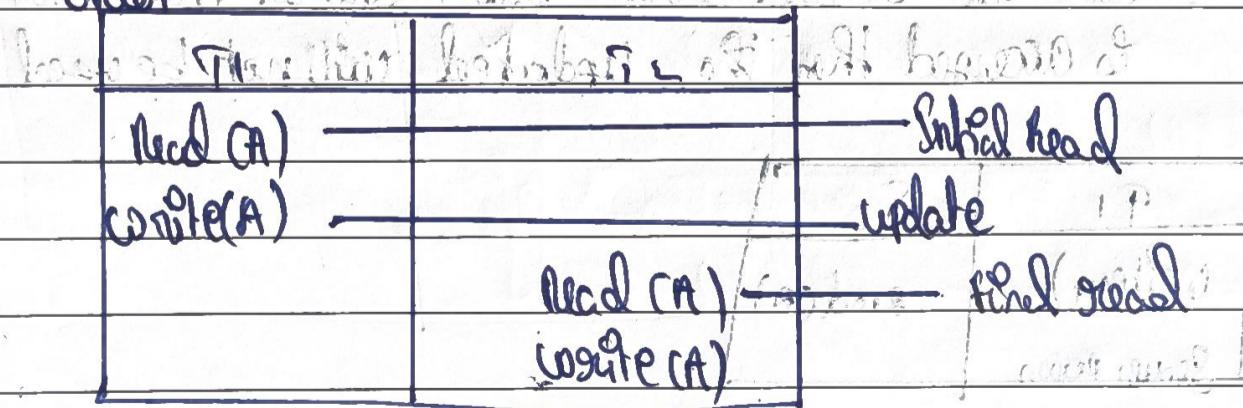
They can be ordered as

T ₁	T ₂	(A, A)	(A, B)
Read(A)			
Write(A)			
Read(A)			
	Read(A)	(A, A)	(A, B)
	Write(A)	(A, B)	

View Serializability:

- * Two schedules S₁ and S₂ are said to be view equivalent if they satisfy the following condition
 - 1) Initial read
 - 2) Update read
 - 3) Final Read

* Simple for each transaction, the Data must be Accessed in same Order and nothing will happen in between at any time.



Concurrency

- * If same data is accessed by different users at the same time, then it leads to Concurrency problems.
- * To control them Concurrency Control is used to fix errors.

Problems in concurrent execution:

1) Lost update problem

T1	T2
R(A)	
W(B)	
R(B)	R(A)
	W(A)
B = A+200	R(B)
	B = A+10
	W(B)

Thus one of the updates will be lost.

2) Dirty Read problem:

- * When a data is updated but not committed due to Server error then if another transaction is accessed then the updated will not be used.

T1	T2
Write(A)	Read(A)
Server down	Client down

3) Unrepeatable read problem:

* when one reads once , and the other updates, next time if the old user reads the data will be different

T1 R(A)	T2 W(A) A=A+100 R(A)

Concurrency protocol

* It ensures ACID properties of transactions.

* The types are

i) Lock Based Concurrency Protocol.

ii) Time Stamp CCP

iii) Validation based CCP.

Lock Based protocol

i) Simple lock protocol

ii) Pre claiming lock protocol

iii) Two phase locking protocol

iv) Strict two phase locking protocol

1) Simple lock protocol : (out of syllabus)

- * It locks all transactions before inserting, updating and deleting and then the transactions are unlocked.

2) Pre-claiming lock protocol : (out of syllabus)

- * It assesses transactions and determines which data elements require lock.
- * The transaction commences if locks are obtained, whenever the transaction is finished the lock is released.

3) Two phase locking protocol: (Syllabus)

It has two phase:

- 1) Growing phase: New locks are acquired on data items but none of these block locks can be released.
- 2) Shrinking phase: The existing locks are released.

Ex	T1	T2	Lock stability
{ lock (A)	-	-	lock - S (A)
lock (B)	-	-	lock - S (B)
{ unlock (A)	-	-	lock - X (C) and unlock (A)
unlock (B)	-	-	lock - X (C) and unlock (B)
			unlock (A) and unlock (B)

Dead Locks:

In two phase locking, it arises when both transaction requires data for each other but they are locked.

T₁

lock A(B)

lock B(C)

deadlock is occurred

Blocking:

If one transaction holds a lock for a long time, other transactions that need the same lock are waiting and blocked.

Time Stamp Ordering Protocol

* It's a method of Concurrency Control

* Transactions are serialized with time stamps

* When an actions are performed the time stamp is stored

Ex: reading and writing with a time stamp

T₁ Read X → RTS → Read TimeStamp

T₁ Read X → RTS(X) → TS(T₁) → TimeStamp

T₂ write X → WTS(X) = TS(T₁)

Rules:

Allows only if time stamp is newer and rejects old time stamp

* If T is older than last write → reject & T (abort T)

If T is older → Reject (Abort T)

New Version Concurrency Control:

When ever a data is updated it does not replace old but rather is a new version of it.

Deadlock Handling

A system is in deadlock state if there exists a set of transactions such that every transaction in the set is waiting for another transaction in the set.

Deadlock prevention

- * For a larger DB, deadlock prevention method is suitable
- * A deadlock can be prevented if the resources are allotted in such a way that deadlock never occurs
- * The DBMS analyzes the operations whether they can create deadlock situation or not, if they do the transaction is never allowed to be executed

Two techniques: i) wait; ii) would wait

i) older transaction needs a data item held by younger transaction - older transaction waits
younger transaction needs a data item held by older transaction - younger transaction dies

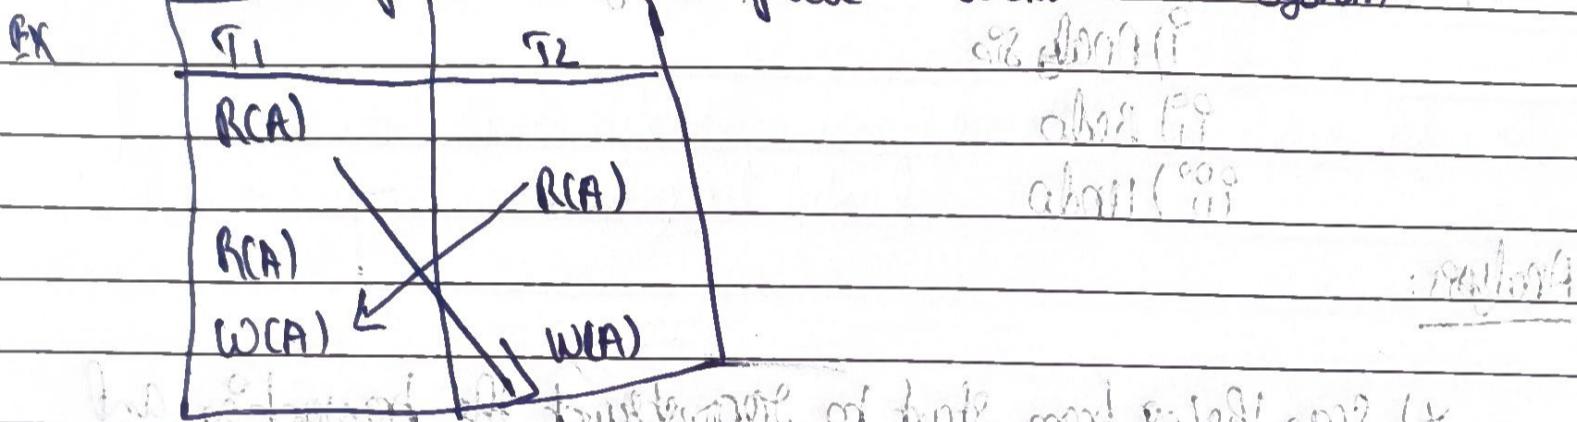
ii) older transaction need a data item held by younger
younger dies

younger transaction needs a data item held by older transaction - younger dies

Deadlock Detection:

- * It is done by wait-for graph method.
- * If deadlock occurrence is detected, then the system must try to recover from it.
- * In this method, a graph is created based on the transaction and their locks.
- * If the created graph has a cycle, then there is a deadlock.
- * The system keeps checking the graph if there is any cycle in the group ($n = \{U, E\}$).

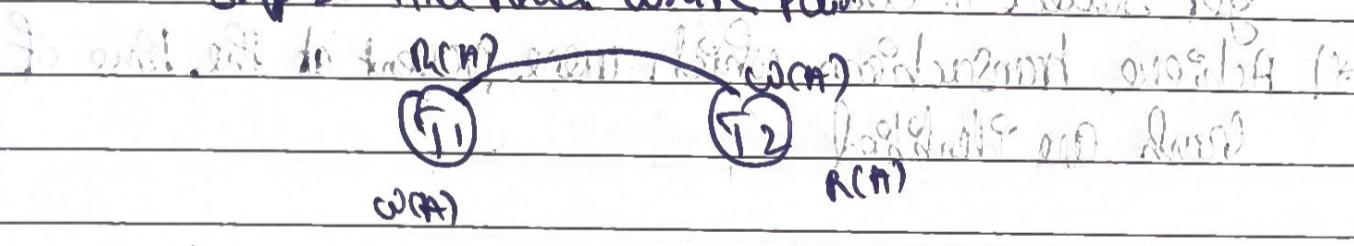
- * A set of vertices consists of all transactions in the system.



Step 1: Draw vertices for all transactions which

can hold Resource and can't make further progress.

Step 2: Find Read write path of Resource.



Step 2: Repeat

Step 3: Since cycle is detected deadlock is found.

Since cycle is detected deadlock is found.

ARIES Algorithm

ARIES → Algorithm for Recovery Isolation Exploiting Semantics

* It is a recovery algorithm

* Algorithm for Recovery Isolation Exploiting Semantics

* It is based on write ahead log protocol (WAL)

* When database crashes during some transaction processing, we have the logs that got saved to the disk

* ARIES has 3 phases that occur in following orders

i) Analyze

ii) Redo

iii) Undo

Analysis:

* Scan the log from start to reconstruct the transaction and dirty page table and mark which pages are dirty.

* Dirty page contain data that has been changed but not yet saved in disk.

* Achieve transactions which were present at the time of crash are identified

Redo:

* It is started only after completing analysis

* The log is read forward and each update is redone

Undo:

* It is started after Redo process

* The log is scanned backward and updates to corresponding active transaction are done

Advantages: It is flexible, concurrency control protocol is supported
Independent recovery of every page.